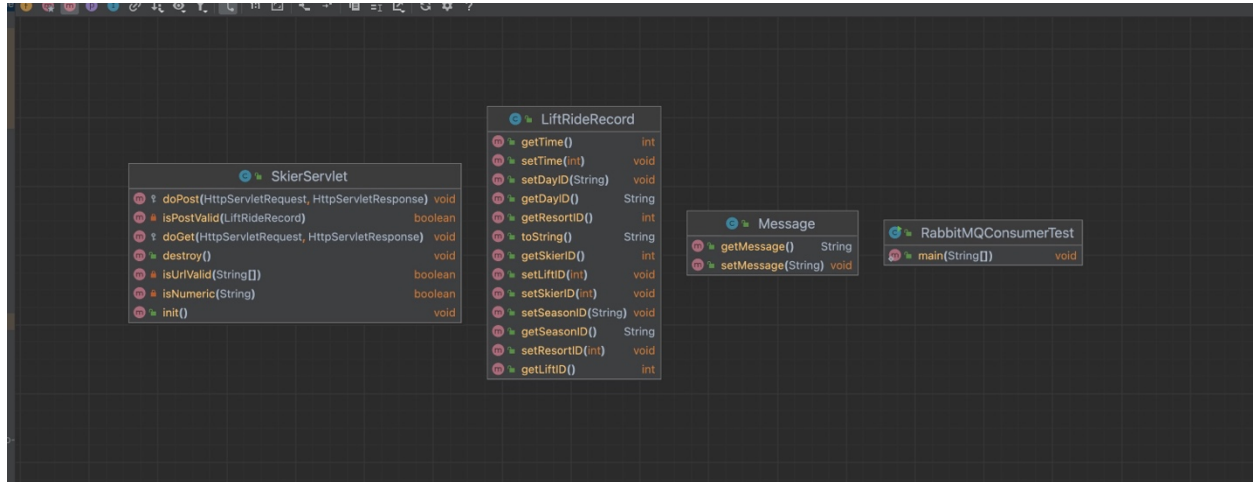


Assignment 3 Zegui Jiang

LiftServer Remain Same with Assignment 2



The server upgraded the functionality for processing messages. Firstly, it connects to RabbitMQ, then it sends the received message queue to RabbitMQ.

`Init()`: Reads and configures settings from the `rabbitmq.conf` configuration file found in resources. It extracts the RabbitMQ queue name and the number of channels from the properties file. All channels are stored in a channel pool, which is a Linked Blocking Queue.

`Destroy()`: Closes all open resources when the process is terminated.

`IsPostValid` and `IsUrlValid`: These functions are likely responsible for validating the incoming request's POST data and the URL, respectively.

`doPost`: Handles incoming POST requests by receiving a message and forwarding it to RabbitMQ.

RabbitMq Lift Record Consumer

<pre> class LiftRideRecord { LiftRideRecord(int, int, int, String, String, int) skierID: int dayID: String seasonID: String resortID: int time: int liftID: int setSeasonID(String): void setSkierID(int): void getTime(): int getSkierID(): int getLiftID(): int getSeasonID(): String getResortID(): int setDayID(String): void getDayID(): String setTime(int): void setLiftID(int): void toString(): String setResortID(int): void </pre>	<pre> class DynamoDBHelper { DynamoDBHelper() successCount: long RETRY_COUNT: int tableName: String dynamoDbClient: DynamoDbClient region: Region BATCH_INSERT_SIZE: int putItemRequests: List<WriteRequest> BACK_OFF_TIME: int checkTableExists(String): boolean getSuccessCount(): long flush(): void batchInsert(Map<String, AttributeValue>): void init(): void insert(Map<String, AttributeValue>): void createLiftRecordTable(): void </pre>	<pre> class DynamoDbQuery { DynamoDbQuery() region: Region SEASON_ID_INDEX: int LIFT_ID_INDEX: int TIME_ID_INDEX: int dynamoDbClient: DynamoDbClient DAY_ID_INDEX: int tableName: String query1(String, String): void query4(String, String): void query3(String): void init(): void query2(String): void main(String[]): void </pre>
<pre> class Consumer { Consumer() liftRecordsMap: Map<Integer, List<JsonObject>> basicQos: int numberOfThread: int connectionFactory: ConnectionFactory rabbitMQName: String properties: Properties checkNumberOfRecordInDDB(ScheduledExecutorService): void main(String[]): void shutDown(ScheduledExecutorService): void </pre>	<pre> class ConsumerThread { ConsumerThread(Connection, String, int) gson: Gson queueName: String basicQos: int connection: Connection run(): void </pre>	

The RabbitMQ consumer, building upon Assignment 2, has integrated DynamoDB as the database to store lift ride records.

DynamoDBHelper

Init: Sets up the DynamoDbClient with the specified AWS region (Region.US_WEST_2). If initialization fails, it prints an error message.

1 - Create table: Checks if the LiftRideRecordTable table exists. If not, it creates the table with a specific schema designed to support queries related to ski lift rides.

This schema includes: Primary key: **skierID** (HASH) and **liftInfo** (RANGE) for identifying records. **skierID**: skier id is partition key because most of query want to check attributes with different skier id

LiftInfo: This is little design in this table. I want to carry most info in to the table. Liftinfo is the concatenate String with symbol “:”, we can decouple when we want these information

Global Secondary Index (GSI): ResortAndDay with resortID (HASH) and dayID (RANGE) to support queries on unique skiers visiting a resort on a specific day. This index includes seasonID

and skierID as non-key attributes. The table and index are provisioned with specific read and write capacity units.

Reason to enable GSI with ResortID and DayID, because the query "How many unique skiers visited resort X on day N?" required scan all over resort id, this will result fast scan.

2 - Inserting Records:

Single inserts (insert method): Inserts a single item into the LiftRideRecordTable. If the insertion fails, it prints an error message along with the item details.

Batch inserts (batch Insert and flush methods): Accumulates insert requests and performs a batch insert when the number of accumulated requests reaches a threshold (BATCH_INSERT_SIZE). The flush method processes the batch insert request; handling retries with exponential backoff (BACK_OFF_TIME) in case of unprocessed items. It keeps track of successfully processed items in success Count.

DynamoDbQuery

Since I could not find a good way to use query in AWS console, I choose use java code to implement these four queries:

"For skier N, how many days have they skied this season?"

"For skier N, what are the vertical totals for each ski day?" (calculate vertical as liftID*10)

"For skier N, show me the lifts they rode on each ski day"

"How many unique skiers visited resort X on day N?"

The result is:

```
/Users/monarch/Library/Java/JavaVirtualMachines/corretto-11.0.22/Contents/Home/bin/java ...  
For Query 1  
For skier 65061, 1 days have skied in season 2024?  
For Query 2  
For Skier: 65061, In Day: 2024:1, Total Vertical: 2030  
For Query 3  
For Skier: 65061, In Day: 2024:1, Lifts Ridden: [12, 2, 15, 26, 4, 17, 39, 29, 19, 8, 32]  
For Query 4  
Number of unique skiers visited resort 10 on day 1: 19910  
  
Process finished with exit code 0
```

Challenge on Dynamo DB deployment and problem solving.

Single insert and batch insert:

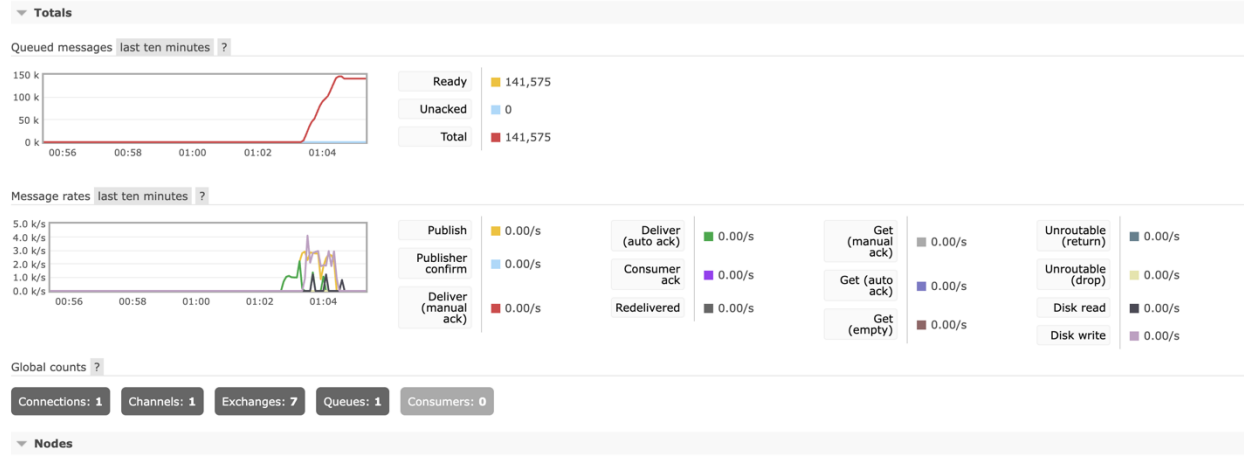
When I try using single insert, the speed of writing to the database becomes very slow, and the latency is very high. Therefore, I used batch write logic, which allows writing multiple records at the same time. Moreover, I implemented retry and back-off time logic to enhance stability.

WCU&RCU:

When setting the writeCapacityUnits and readCapacityUnits, I tried many combinations, but none worked well. There would be a lot of backlog in RabbitMQ, and the writing of consumers would also be blocked. Later, I adjusted the WCU to on-demand. DynamoDB will automatically adjust the size and speed of write and read, preventing blockages. RabbitMQ will also be healthier.

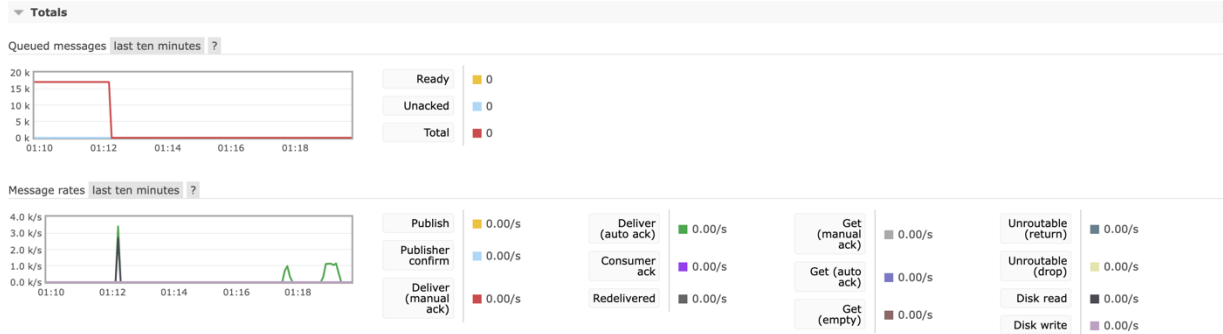
Before:

Overview



After:

Overview



Deployment Topology:

Client -> TomcatServer -> RabbitMq <- Consumer -> DynamoDB

Instances (7) [Info](#)

[Refresh](#) [Connect](#) [Instance state](#) [Actions](#) [Launch instances](#)

<input type="checkbox"/>	Name ↗	Instance ID	Instance state ↕	Instance type ↕	Status check	Alarm status	Availability Zone ↕	Public IPv4 DNS ↕	Public IPv4 ... ↕	Elastic IP
<input type="checkbox"/>	Consumer	i-05a9c6ee256d05002	Running 🔍	t2.micro	2/2 checks passed	View alarms +	us-west-2a	ec2-34-212-147-15.us-...	34.212.147.15	-
<input type="checkbox"/>	TomcatServer1	i-095e7465ee7d42b12	Running 🔍	t2.micro	2/2 checks passed	View alarms +	us-west-2b	ec2-34-210-77-221.us-...	34.210.77.221	-
<input type="checkbox"/>	RabbitMQServer	i-0c070b5928025afb7	Running 🔍	t2.micro	2/2 checks passed	View alarms +	us-west-2a	ec2-44-237-68-220.us-...	44.237.68.220	44.237.68.220
<input type="checkbox"/>	LiftRideConsumer	i-0d81fe16da8ab60cc	Stopped 🔍	t2.micro	-	View alarms +	us-west-2a	-	-	-
<input type="checkbox"/>	6650LabWebServer	i-046718c5b3fb12339	Running 🔍	t2.micro	2/2 checks passed	View alarms +	us-west-2a	ec2-35-165-79-236.us-...	35.165.79.236	-
<input type="checkbox"/>	TomcatServer3	i-059f0772610b8933b	Stopped 🔍	t2.micro	-	View alarms +	us-west-2c	-	-	-
<input type="checkbox"/>	TomcatServer2	i-037e6630a80598b81	Stopped 🔍	t2.micro	-	View alarms +	us-west-2c	-	-	-

Saving log:

```
Successful saved 7800 item in to Dynamo DB
Successful saved 7825 item in to Dynamo DB
Successful saved 7850 item in to Dynamo DB
Successful saved 7875 item in to Dynamo DB
Successful saved 7900 item in to Dynamo DB
Successful saved 7925 item in to Dynamo DB
Successful saved 7950 item in to Dynamo DB
Successful saved 7975 item in to Dynamo DB
Successful saved 8000 item in to Dynamo DB
Successful saved 8025 item in to Dynamo DB
Successful saved 8050 item in to Dynamo DB
Successful saved 8075 item in to Dynamo DB
```

Saving result:

Items summary

Get live item count

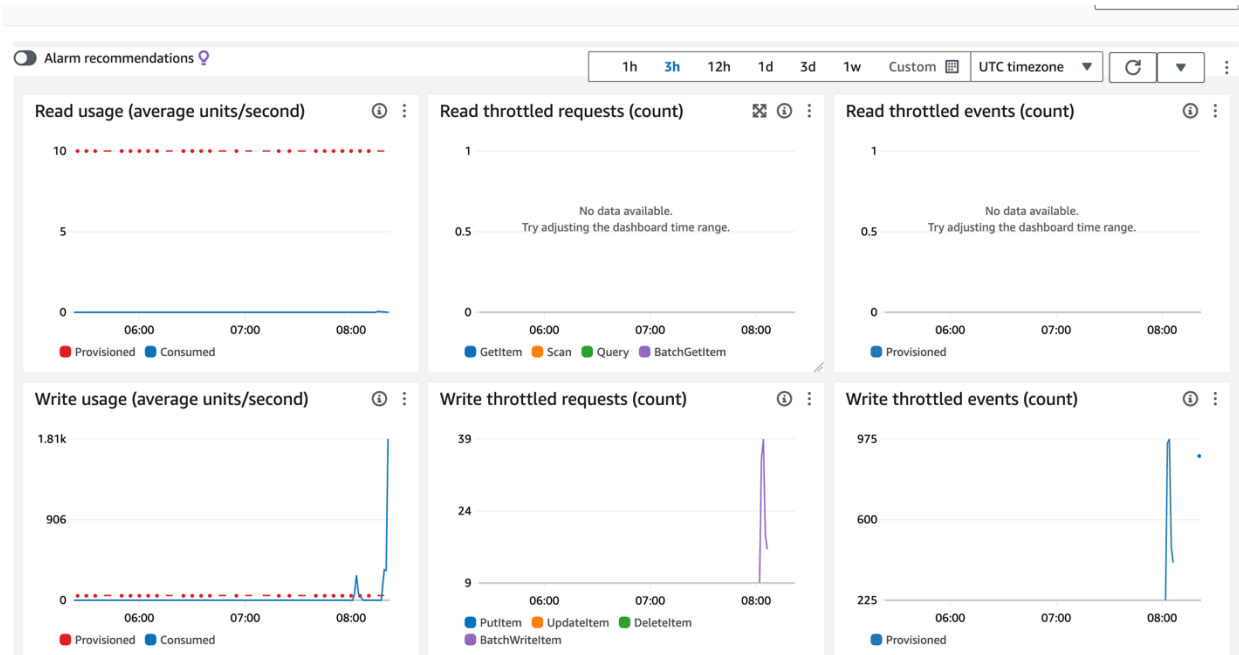
When you choose "Start scan," you will perform a DynamoDB scan to determine the most-recent item count. This scan might consume additional table read capacity units.

⚠ It is not recommended to perform this action on very large tables or tables that serve critical production traffic. You can pause the action at any time to avoid consuming extra read capacity.

Item count	Scan status	Last updated
200,000	✔ Complete	March 28, 2024 01:45:04

Scan again

Cancel



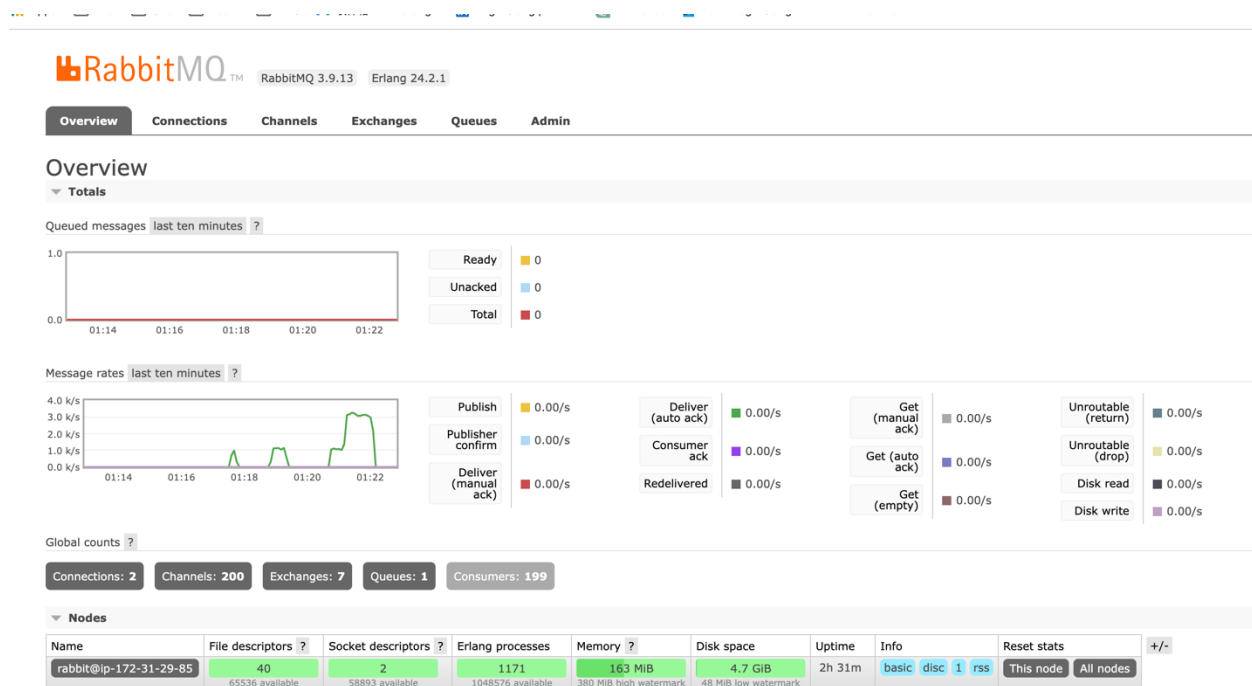
Client test run and RMQ console screen shots showing your best throughput with ideally short, stable queue lengths.

Number of Thread in Client mode 2: 100

Number of Thread in Consumer: 400

```
MultiThreadCall x
/Users/monarch/Library/Java/JavaVirtualMachines/corretto-11.0.22/Contents/Home/bin/java ...
Statistic Metrics
Mean Response Time: 31.749735 ms
Median Response Time: 31.0 ms
P99 Response Time: 58 ms
Min Response Time: 14 ms
Max Response Time: 447 ms
Summary:
Number of thread in process 2: 100
Number of successful requests: 200000
Number of fail requests: 0
Total run time: 82723
Response Time: 0.413615 ms/request
RPS: 2417 requests/second

Process finished with exit code 0
```



Number of Thread in Client mode 2: 200
Number of Thread in Consumer: 400

Statistic Metrics

Mean Response Time: 32.28399604075226 ms

Median Response Time: 31.0 ms

P99 Response Time: 60 ms

Min Response Time: 16 ms

Max Response Time: 265 ms

Summary:

Number of thread in process 2: 200

Number of successful requests: 200000

Number of fail requests: 0

Total run time: 55686

Response Time: 0.27843 ms/request

RPS: 3591 requests/second

Process finished with exit code 0



Overview Connections Channels Exchanges Queues Admin

Overview

Totals

Queued messages last ten minutes ?



Ready 0
Unacked 0
Total 0

Message rates last ten minutes ?



Publish 0.00/s
Publisher confirm 0.00/s
Deliver (manual ack) 0.00/s
Deliver (auto ack) 0.00/s
Consumer ack 0.00/s
Redelivered 0.00/s
Get (manual ack) 0.00/s
Get (auto ack) 0.00/s
Get (empty) 0.00/s
Unroutable (return) 0.00/s
Unroutable (drop) 0.00/s
Disk read 0.00/s
Disk write 0.00/s

Global counts ?

Connections: 2 Channels: 401 Exchanges: 7 Queues: 1 Consumers: 400

Nodes

Name	File descriptors ?	Socket descriptors ?	Erlang processes	Memory ?	Disk space	Uptime	Info	Reset stats	+/-
rabbit@ip-172-31-29-85	40 65536 available	2 58893 available	1975 1048576 available	183 MiB 380 MiB high watermark 48 MiB low watermark	4.7 GiB	2h 58m	basic disc 1 rss	This node All nodes	

Churn statistics

Ports and contexts

Export definitions

Import definitions

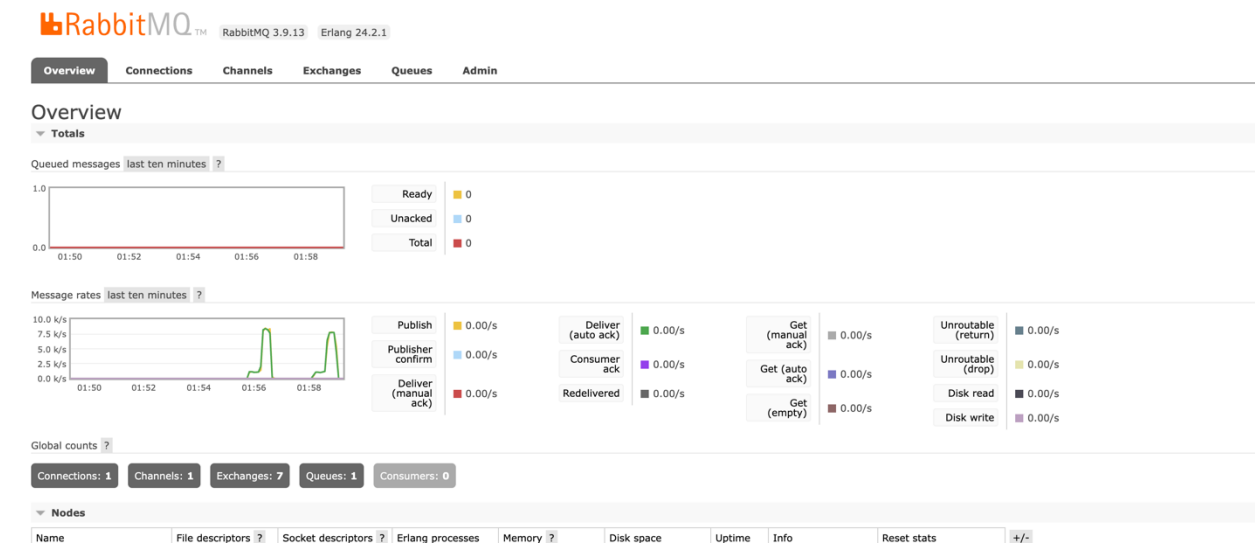
HTTP API Server Docs Tutorials Community Support Community Slack Commercial Support Plugins GitHub Changelog

Number of Thread in Client mode 2: 400

Number of Thread in Consumer: 400

```
MultiThreadCall x
/Users/monarch/Library/Java/JavaVirtualMachines/corretto-11.0.22/Contents/Home/bin/java ...
Statistic Metrics
Mean Response Time: 48.05203 ms
Median Response Time: 48.0 ms
P99 Response Time: 99 ms
Min Response Time: 15 ms
Max Response Time: 274 ms
Summary:
Number of thread in process 2: 400
Number of successful requests: 200000
Number of fail requests: 0
Total run time: 48896
Response Time: 0.24448 ms/request
RPS: 4090 requests/second

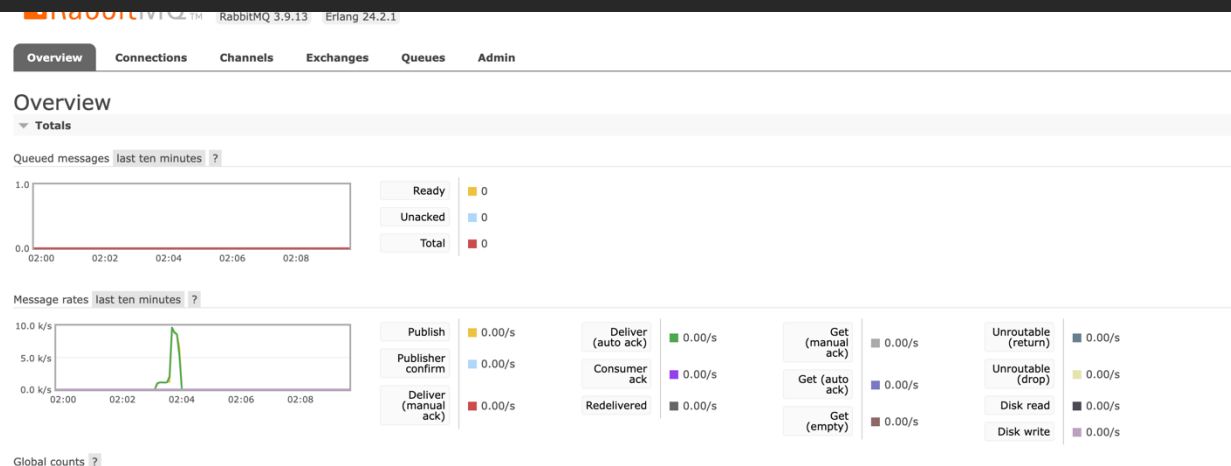
Process finished with exit code 0
```



Number of Thread in Client mode 2: 800

Number of Thread in Consumer: 400

```
Statistic Metrics
Mean Response Time: 77.16995502248875 ms
Median Response Time: 67.0 ms
P99 Response Time: 202 ms
Min Response Time: 15 ms
Max Response Time: 770 ms
Summary:
Number of thread in process 2: 800
Number of successful requests: 200000
Number of fail requests: 0
Total run time: 42954
Response Time: 0.21477 ms/request
RPS: 4656 requests/second
```



Number of Thread in Client mode 2: 1000


Number of Thread in Consumer: 400

```
Statistic Metrics
Mean Response Time: 93.34594989029443 ms
Median Response Time: 84.0 ms
P99 Response Time: 247 ms
Min Response Time: 14 ms
Max Response Time: 1293 ms
Summary:
Number of thread in process 2: 1000
Number of successful requests: 200000
Number of fail requests: 0
Total run time: 44324
Response Time: 0.22162 ms/request
RPS: 4512 requests/second
```

Overview

Totals

Queued messages last ten minutes ?



Ready	0
Unacked	0
Total	0

Message rates last ten minutes ?



Publish	0.00/s
Publisher confirm	0.00/s
Deliver (manual ack)	0.00/s

Deliver (auto ack)	0.00/s
Consumer ack	0.00/s
Redelivered	0.00/s

Get (manual ack)	0.00/s
Get (auto ack)	0.00/s
Get (empty)	0.00/s

Unroutable (return)	0.00/s
Unroutable (drop)	0.00/s
Disk read	0.00/s
Disk write	0.00/s

Global counts ?

Connections: 1 Channels: 1 Exchanges: 7 Queues: 1 Consumers: 0